

Data Wrangling

Billy Lozowski

2025-03-21

Contents

1. Download two .csv files from Canvas called DiversityData.csv and Metadata.csv, and read them into R using relative file paths.

```
diversity.data <- read.csv("data/DiversityData.csv")
str(diversity.data)
```

```
## 'data.frame':    70 obs. of  5 variables:
## $ Code          : chr  "S01_13" "S02_16" "S03_19" "S04_22" ...
## $ shannon       : num  6.62 6.61 6.66 6.66 6.61 ...
## $ invsimpson    : num  211 207 213 205 200 ...
## $ simpson       : num  0.995 0.995 0.995 0.995 0.995 ...
## $ richness      : int  3319 3079 3935 3922 3196 3481 3250 3170 3657 3177 ...
```

```
meta.data <- read.csv("data/Metadata.csv")
str(meta.data)
```

```
## 'data.frame':    70 obs. of  5 variables:
## $ Code          : chr  "S01_13" "S02_16" "S03_19" "S04_22" ...
## $ Crop          : chr  "Soil" "Soil" "Soil" "Soil" ...
## $ Time_Point    : int   0 0 0 0 0 0 6 6 6 6 ...
## $ Replicate     : int   1 2 3 4 5 6 1 2 3 4 ...
## $ Water_Imbided: chr   "na" "na" "na" "na" ...
```

```
library(tidyverse)
```

```
## Warning: package 'stringr' was built under R version 4.4.2
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

2. Join the two dataframes together by the common column 'Code'. Name the resulting dataframe alpha.

```
alpha <- left_join(meta.data, diversity.data, by = "Code")
```

3. Calculate Pielou's evenness index: Pielou's evenness is an ecological parameter calculated by the Shannon diversity index (column Shannon) divided by the log of the richness column.
 - a. Using mutate, create a new column to calculate Pielou's evenness index.
 - b. Name the resulting dataframe alpha_even.

```
alpha_even <- alpha %>%  
  mutate(Pielou_evenness = shannon / log(richness))
```

4. Using tidyverse language of functions and the pipe, use the summarise function and tell me the mean and standard error evenness grouped by crop over time.
 - a. Start with the alpha_even dataframe
 - b. Group the data: group the data by Crop and Time_Point.
 - c. Summarize the data: Calculate the mean, count, standard deviation, and standard error for the even variable within each group.
 - d. Name the resulting dataframe alpha_average

```
alpha_average <- alpha_even %>%  
  group_by(Crop, Time_Point) %>%  
  summarise(Mean_evenness = mean(Pielou_evenness),  
            n = n(),  
            sd = sd(Pielou_evenness)) %>%  
  mutate(se = sd/sqrt(n))
```

```
## 'summarise()' has grouped output by 'Crop'. You can override using the  
## '.groups' argument.
```

5. Calculate the difference between the soybean column, the soil column, and the difference between the cotton column and the soil column
 - a. Start with the alpha_average dataframe
 - b. Select relevant columns: select the columns Time_Point, Crop, and mean.even.
 - c. Reshape the data: Use the pivot_wider function to transform the data from long to wide format, creating new columns for each Crop with values from mean.even.
 - d. Calculate differences: Create new columns named diff.cotton.even and diff.soybean.even by calculating the difference between Soil and Cotton, and Soil and Soybean, respectively.
 - e. Name the resulting dataframe alpha_average2

```
alpha_average2 <- alpha_average %>%  
  select(Time_Point, Crop, Mean_evenness) %>%  
  pivot_wider(names_from = Crop, values_from = Mean_evenness) %>%  
  mutate(diff.cotton.even = Soil - Cotton,  
         diff.soybean.even = Soil - Soybean)
```


6. Connecting it to plots

- a. Start with the `alpha_average2` dataframe
- b. Select relevant columns: select the columns `Time_Point`, `diff.cotton.even`, and `diff.soybean.even`.
- c. Reshape the data: Use the `pivot_longer` function to transform the data from wide to long format, creating a new column named `diff` that contains the values from `diff.cotton.even` and `diff.soybean.even`.
 - i. This might be challenging, so I'll give you a break. The code is below.

`pivot_longer(c(diff.cotton.even, diff.soybean.even), names_to = "diff")` *I had to modify this to get it to run*

- d. Create the plot: Use `ggplot` and `geom_line()` with '`Time_Point`' on the x-axis, the column '`values`' on the y-axis, and different colors for each '`diff`' category. The column named '`values`' come from the `pivot_longer`.

```
alpha_average2 %>%
  select(Time_Point, diff.cotton.even, diff.soybean.even) %>%
  pivot_longer(cols = c(diff.cotton.even, diff.soybean.even),
               names_to = "diff",
               values_to = "values") %>%
  ggplot(aes(x = Time_Point, y = values, color = diff)) +
  geom_line() +
  theme_minimal() +
  labs(x = "Time Point",
       y = "Difference")
```



7. Commit and push a `gfm .md` file to GitHub inside a directory called Coding Challenge 5. Provide me a link to your github written as a clickable link in your `.pdf` or `.docx`

GitHub Link